AI TRUST CASE STUDIES

Case Study 4: ChatGPT Hallucinations

Author

HANDE KRISHNA

GitHub Repository

https://github.com/krishna111-31

RESEARCH INTERESTS

Trustworthy AI, Human–AI Interaction (HAI), Responsible AI, Explainable AI (XAI), Human-Centered AI (HCAI), AI Ethics.

PURPOSE

This document is part of the AI Trust Case Studies project. It provides an educational analysis of a real-world AI system using widely accepted Trustworthy AI principles.

DISCLAIMER

This material is intended for educational and research purposes only. It is based on publicly available information from official reports, peer-reviewed literature, and reputable sources. The analysis reflects the cited evidence and does not represent the views of any organization mentioned. Readers should consult the original references for the most authoritative and up-to-date information.

CASE STUDY 4: CHATGPT HALLUCINATIONS

Topic: Reliability, Explainability, Human Oversight, and Trustworthy AI

LEARNING OBJECTIVES

After studying this case, you will be able to:

- Understand what AI hallucinations are and why they occur.
- Explain the importance of reliability in generative AI systems.
- Analyze the role of human oversight when using Large Language Models (LLMs).
- Identify the Trustworthy AI principles relevant to generative AI.
- Evaluate strategies to reduce hallucinations and improve AI reliability.

1. OVERVIEW

Large Language Models (LLMs), such as ChatGPT, have transformed how people interact with Artificial Intelligence. They can generate human-like text, answer questions, summarize documents, write code, and assist in many professional and educational tasks.

However, one well-known limitation is AI hallucination—the generation of responses that sound plausible but are factually incorrect, fabricated, or unsupported by evidence. Hallucinations can occur because language models predict the most likely next words based on patterns in training data rather than verifying facts against a trusted knowledge source.

Although ChatGPT is useful in many applications, OpenAI and other AI developers recommend that users verify important information, particularly in domains such as healthcare, law, finance, and scientific research.

This case study examines hallucinations from the perspective of Trustworthy AI, emphasizing reliability, explainability, human oversight, and responsible AI use.

2. BACKGROUND

Generative AI systems are trained on large datasets containing books, articles, websites, and other publicly available text. During training, these models learn statistical patterns in language rather than storing or reasoning about facts in the way humans do.

As a result, when the model encounters ambiguous, incomplete, or unfamiliar prompts, it may generate confident but inaccurate information. These incorrect outputs are commonly referred to as hallucinations.

Hallucinations are an active area of AI research, and developers continue to improve model accuracy through techniques such as reinforcement learning, retrieval-augmented generation (RAG), better evaluation methods, and human feedback.

3. TIMELINE

YEAR	EVENT
2022	ChatGPT released by OpenAI for public use.
2023	Rapid adoption in education, business, and software development.
2023	Researchers and organizations report hallucination-related challenges.
2024	Continued improvements in model reliability and evaluation methods.
Present	Hallucination reduction remains an active research area.

4. PROBLEM STATEMENT

Generative AI systems aim to:

- Provide accurate information.
- Answer complex questions.
- Generate useful content.
- Support human decision-making.

The challenge is ensuring that AI-generated information is reliable, especially when users may assume that confident responses are always correct.

5. HOW THE AI SYSTEM WORKED

The process can be summarized as follows:

- User submits a prompt.
- The language model processes the input.
- The model predicts the most probable sequence of words.
- A response is generated.
- The user reviews and verifies the output.

Unlike traditional search engines, language models generate responses rather than retrieving verified facts directly.

6. WHAT WENT WRONG?

Hallucinations occur when the model generates incorrect, fabricated, or unsupported information.

Possible causes include:

- Ambiguous user prompts.

- Limited or conflicting training information.
- Prediction-based text generation.
- Lack of access to verified real-time information in some contexts.

Hallucinations may include:

- Incorrect citations.
- Invented references.
- Fabricated quotations.
- Incorrect factual statements.
- Non-existent legal cases or research papers.

Importantly, hallucinations do not indicate intentional deception. They are a limitation of current generative AI technology.

7. TRUSTWORTHY AI PRINCIPLES INVOLVED

Reliability
AI systems should produce dependable and consistent outputs.
Lesson: Users should verify important information before relying on AI-generated content.

Explainability
Users benefit from understanding the limitations of AI-generated responses.
Lesson: Organizations should clearly communicate AI capabilities and limitations.

Human Oversight
Humans remain responsible for reviewing AI-generated information.
Lesson: AI should support—not replace—human judgment.

Accountability
Organizations deploying AI remain accountable for how the technology is used.
Lesson: Developers should continuously evaluate and improve AI systems.

Governance
Continuous monitoring and evaluation help improve AI quality over time.
Lesson: Responsible governance includes regular testing, benchmarking, and updates.

8. HUMAN–AI INTERACTION PERSPECTIVE

Human–AI Interaction focuses on effective collaboration between people and AI systems.

For generative AI:

- Users should critically evaluate responses.
- AI should communicate uncertainty where appropriate.

- Human expertise remains essential.
- AI should assist decision-making rather than replace expert judgment.

Appropriate trust is achieved when users understand both the strengths and limitations of AI systems.

9. IMPACT

On Users

- Increased productivity.
- Improved access to information.
- Need for careful fact-checking.

On Organizations

- Adoption of AI governance policies.
- Greater emphasis on responsible AI deployment.
- Increased investment in AI evaluation and monitoring.

On AI Research

- New research on hallucination detection.
- Improved retrieval methods.
- Better evaluation benchmarks.
- Enhanced safety mechanisms.

10. LESSONS LEARNED

- AI-generated content should be verified.
- Reliability is essential for trustworthy AI.
- Human oversight remains necessary.
- Explainability improves user trust.
- Continuous evaluation strengthens AI systems.
- Responsible AI requires transparency and governance.

11. BEST PRACTICES

Organizations using generative AI should:

- Verify AI-generated information.
- Use retrieval-based approaches where appropriate.
- Maintain human review for high-impact decisions.

• Educate users about AI limitations.

• Monitor AI performance continuously.

• Document known limitations.

• Follow recognized AI governance frameworks.

12. DISCUSSION QUESTIONS

• Why do AI hallucinations occur?

• How can users reduce the risks associated with hallucinations?

• Why is human oversight important when using generative AI?

• Which Trustworthy AI principles are most relevant?

• Should AI-generated information always be independently verified?

13. KEY TAKEAWAYS

• Hallucinations are a known limitation of current generative AI systems.

• Reliability is a core principle of Trustworthy AI.

• Human oversight remains essential.

• Explainability improves appropriate trust.

• Continuous research aims to reduce hallucinations.

14. CONNECTION TO PREVIOUS CHAPTERS

CHAPTER	CONNECTION
Trustworthy AI	Demonstrates the importance of reliability and accountability.
Human–AI Interaction	Shows AI as a collaborative tool requiring human judgment.
Explainable AI	Highlights the need to communicate AI limitations.
Responsible AI	Reinforces governance and continuous evaluation.

15. REFERENCES

• OpenAI. ChatGPT Safety and Usage Documentation. https://openai.com/

• NIST. (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). https://www.nist.gov/itl/ai-risk-management-framework

• OECD. (2019). OECD AI Principles. https://oecd.ai/

• UNESCO. (2021). Recommendation on the Ethics of Artificial Intelligence. https://www.unesco.org/

• European Commission. (2019). Ethics Guidelines for Trustworthy AI. https://digital-strategy.ec.europa.eu/

• OpenAI. Research and Safety Publications. https://openai.com/research